



COMP3005: Computer Vision

Machine learning for pattern recognition

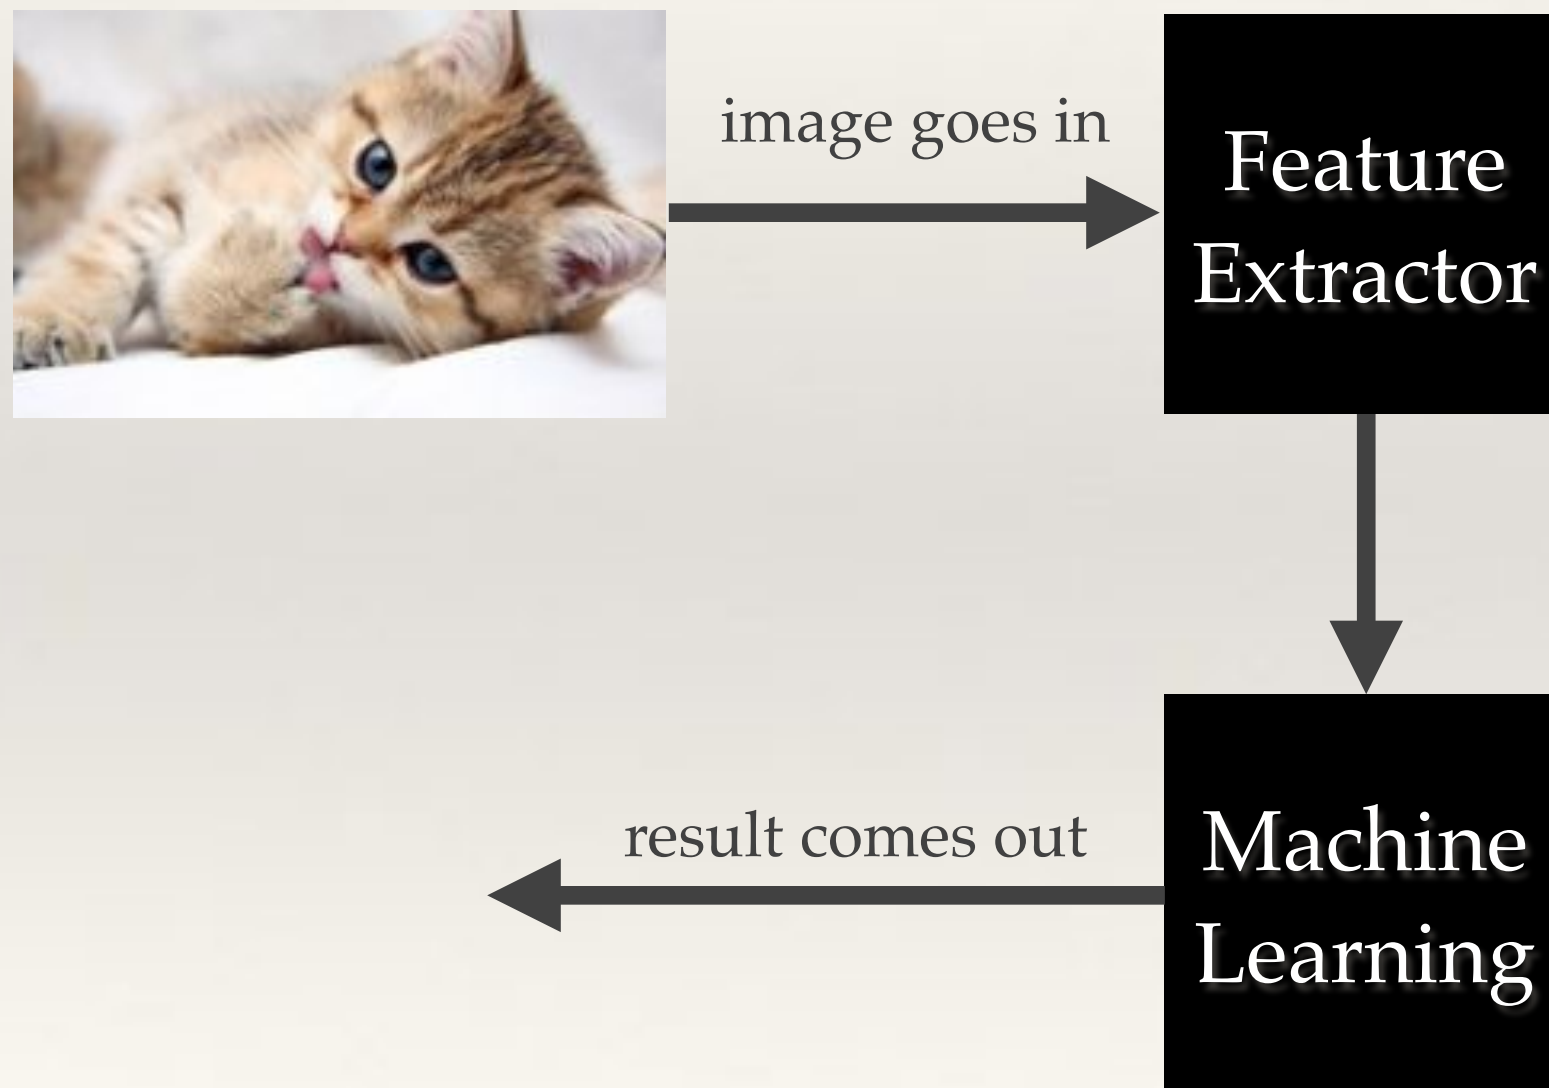
Jonathon Hare
jsh2@ecs.soton.ac.uk

RETURN TO D-STATION.

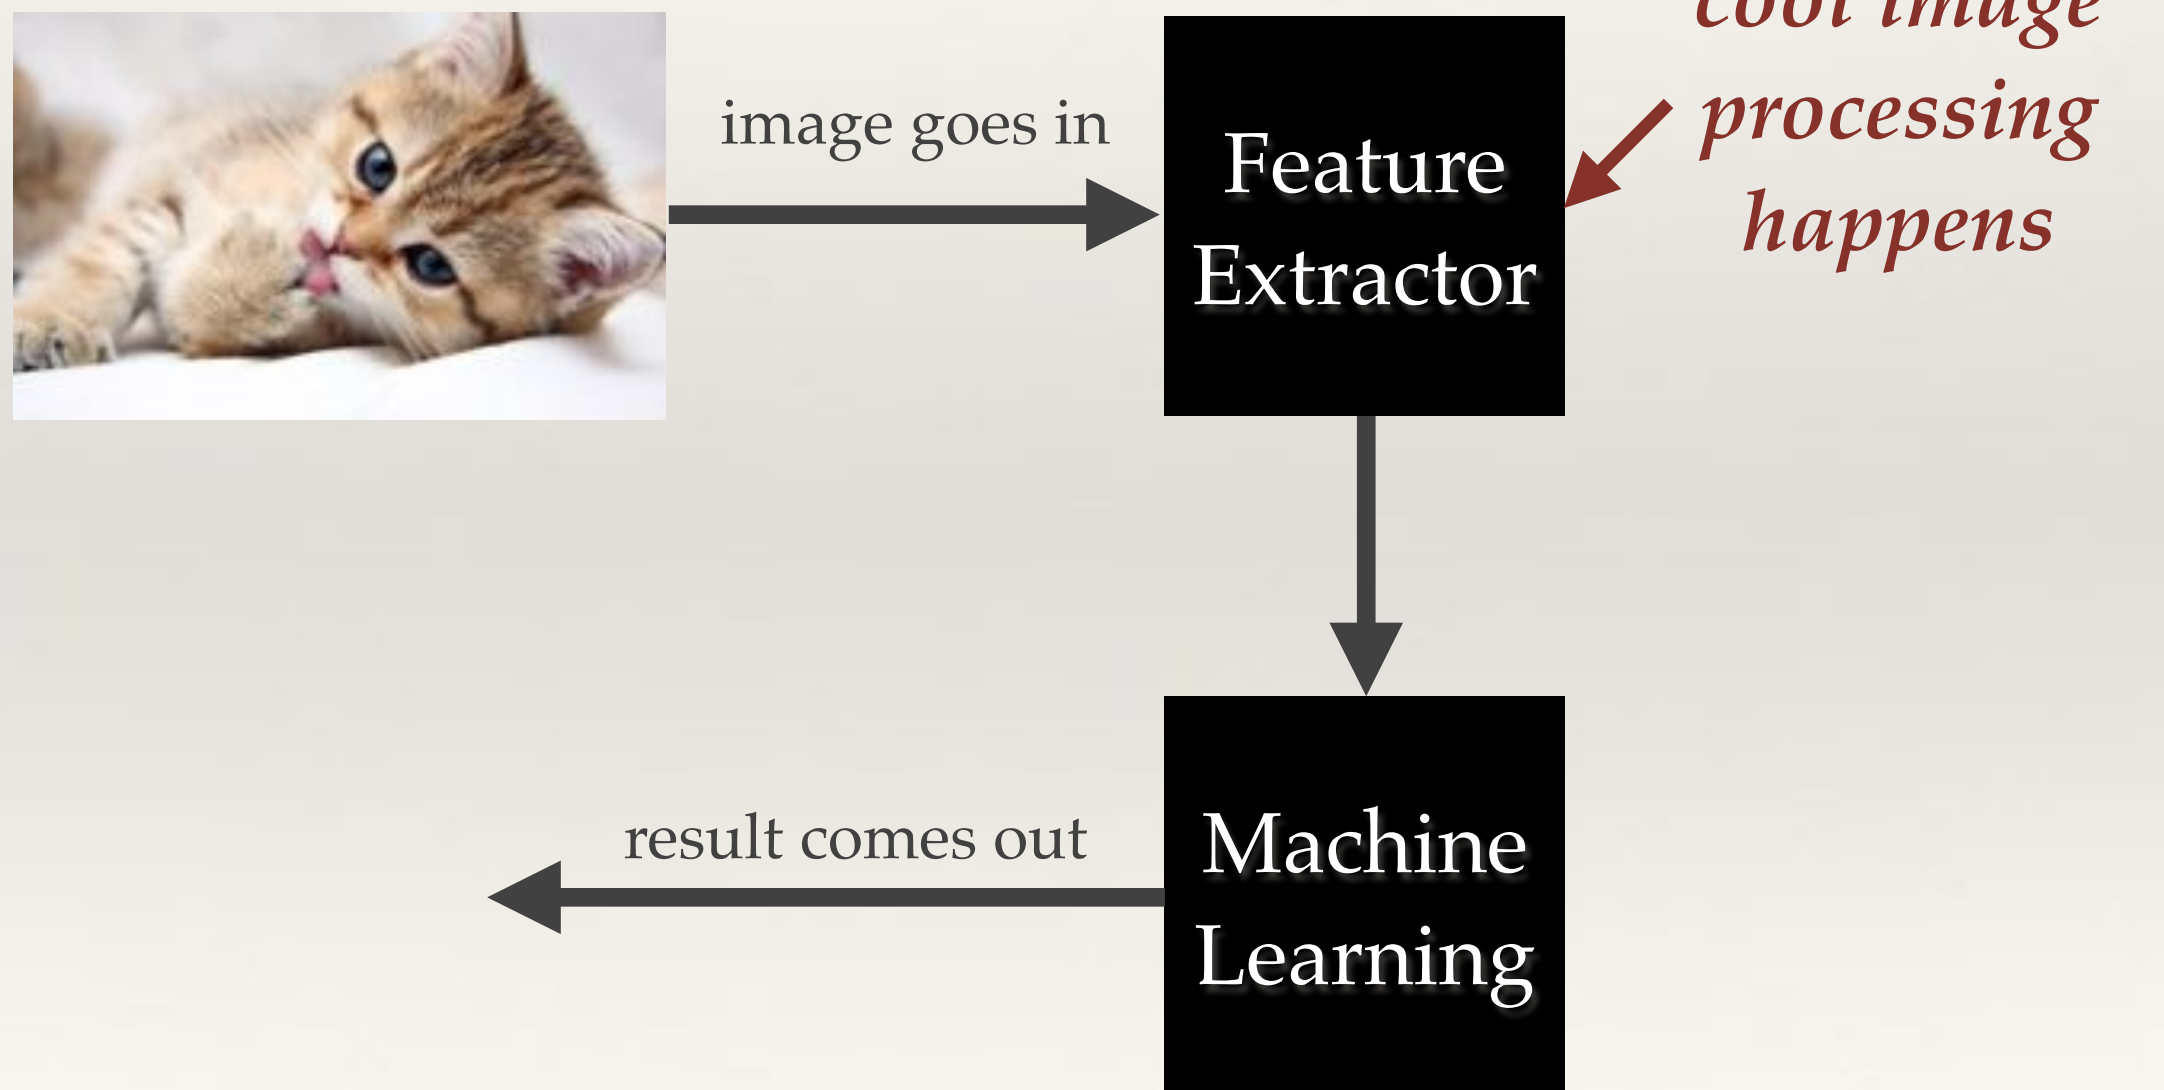
- ❖ Recognising patterns is a large part of computer vision
 - ❖ i.e. recognising text, people, objects, ...
- ❖ Obviously there's a lot of overlap with intelligent algorithms, machine learning and AI.
- ❖ This lecture will cover (recap?) some of the fundamentals of machine learning and introduce how you connect arrays of pixels to machine learning algorithms.

Feature spaces

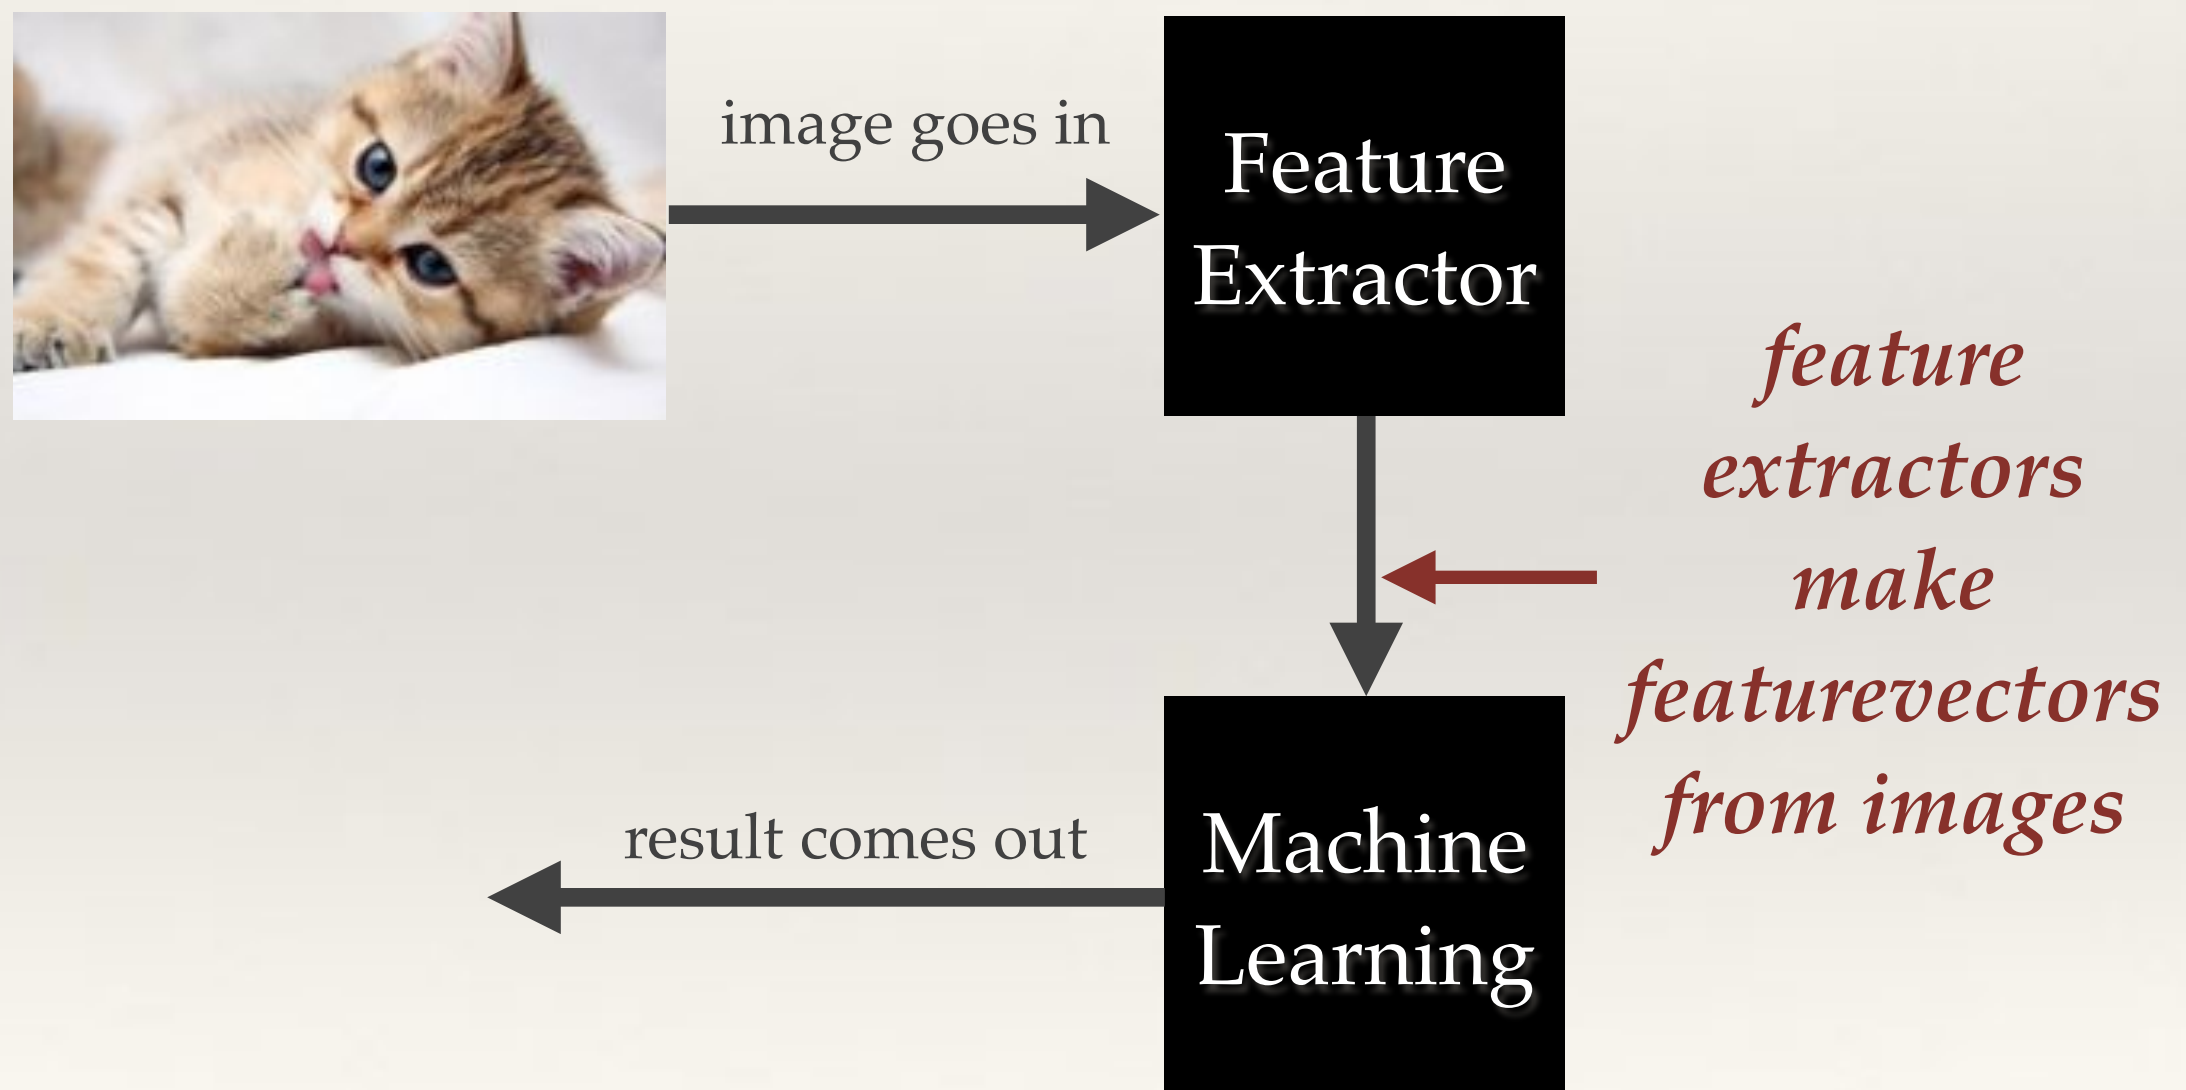
Many computer vision applications involving machine learning take the following form:



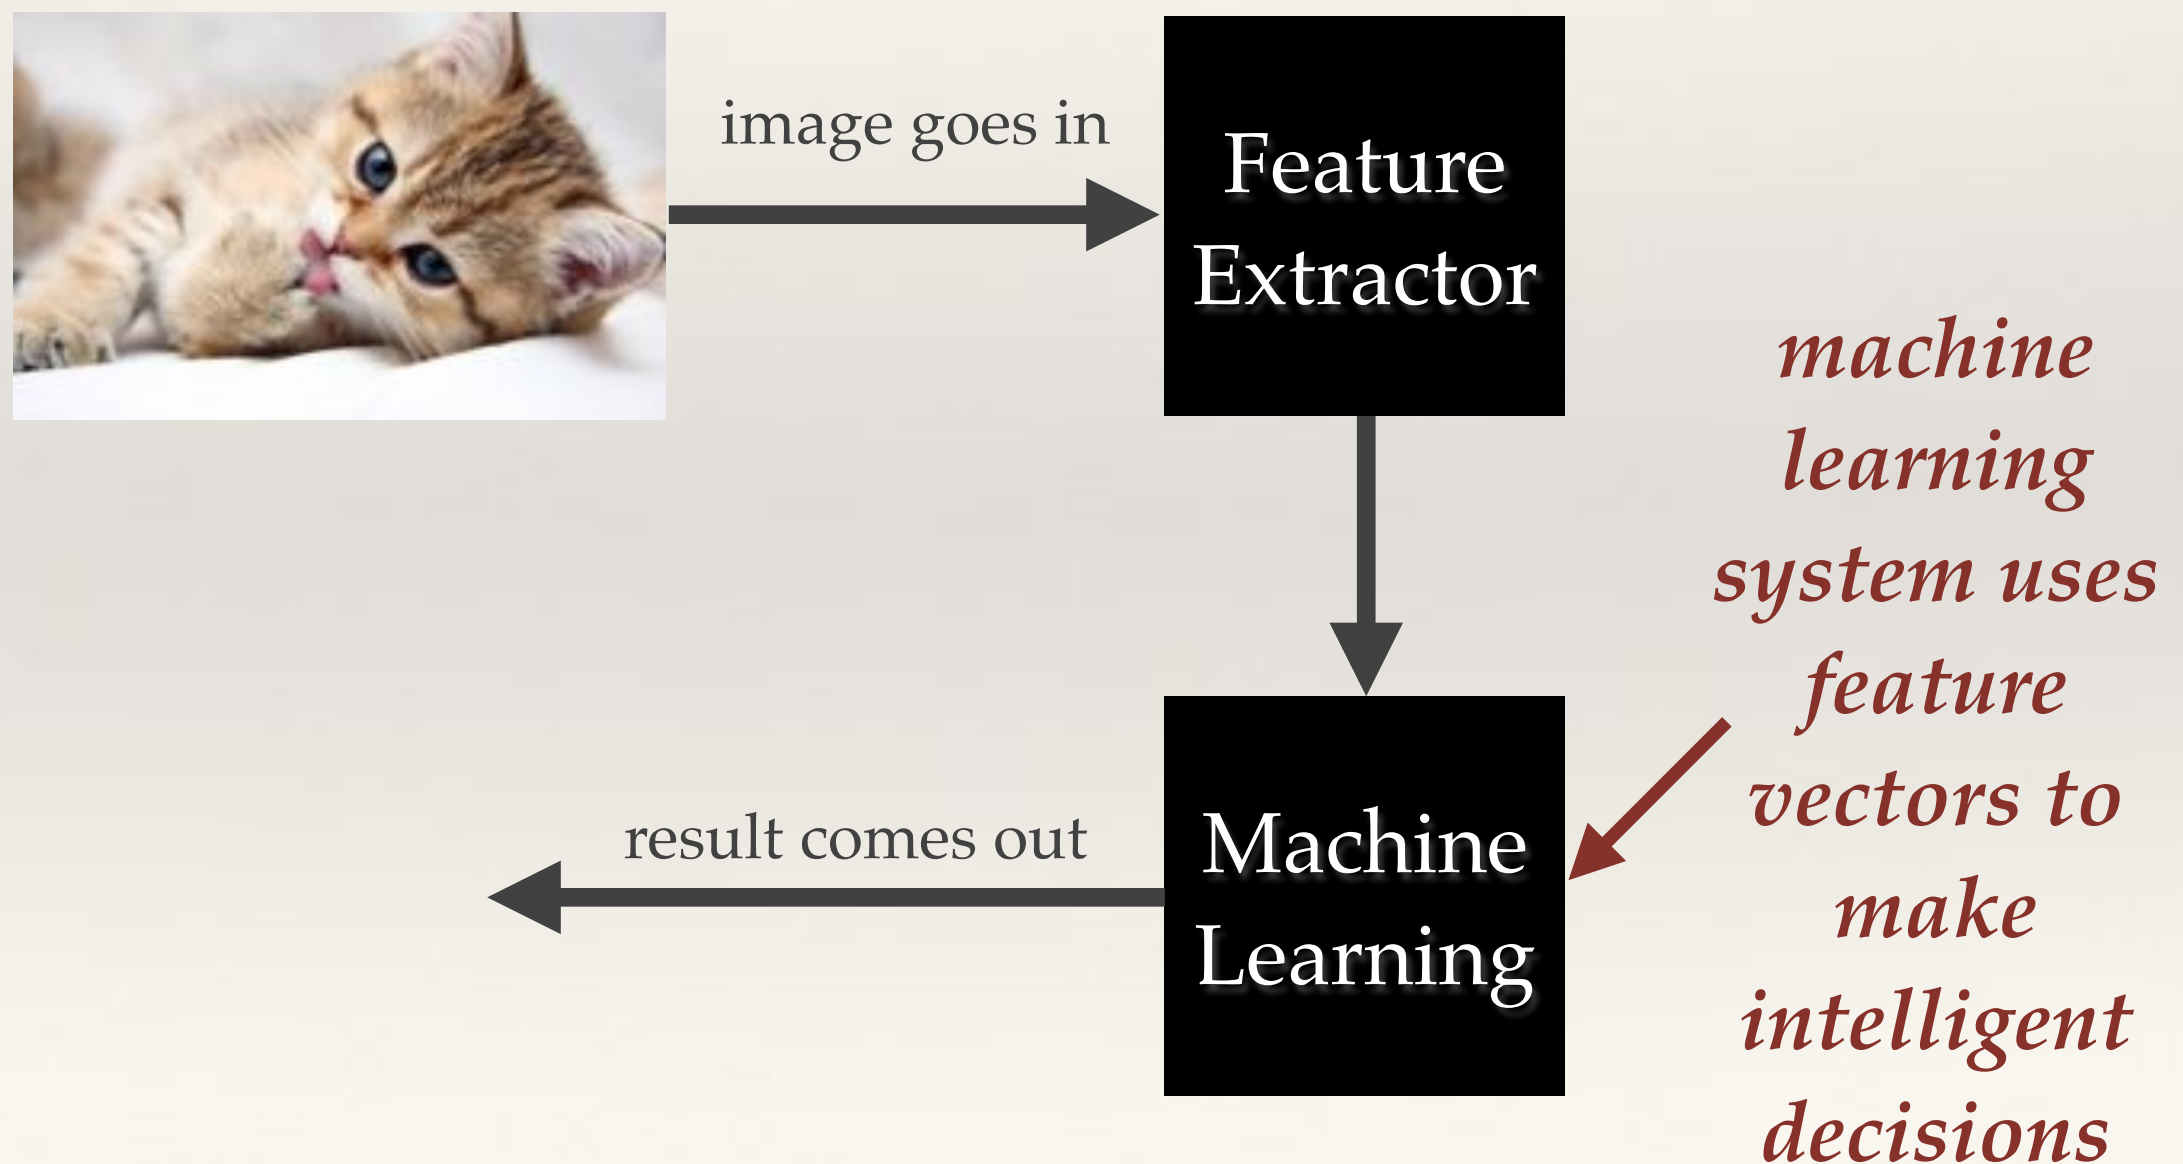
Many computer vision applications involving machine learning take the following form:



Many computer vision applications involving machine learning take the following form:



Many computer vision applications involving machine learning take the following form:



Key terminology

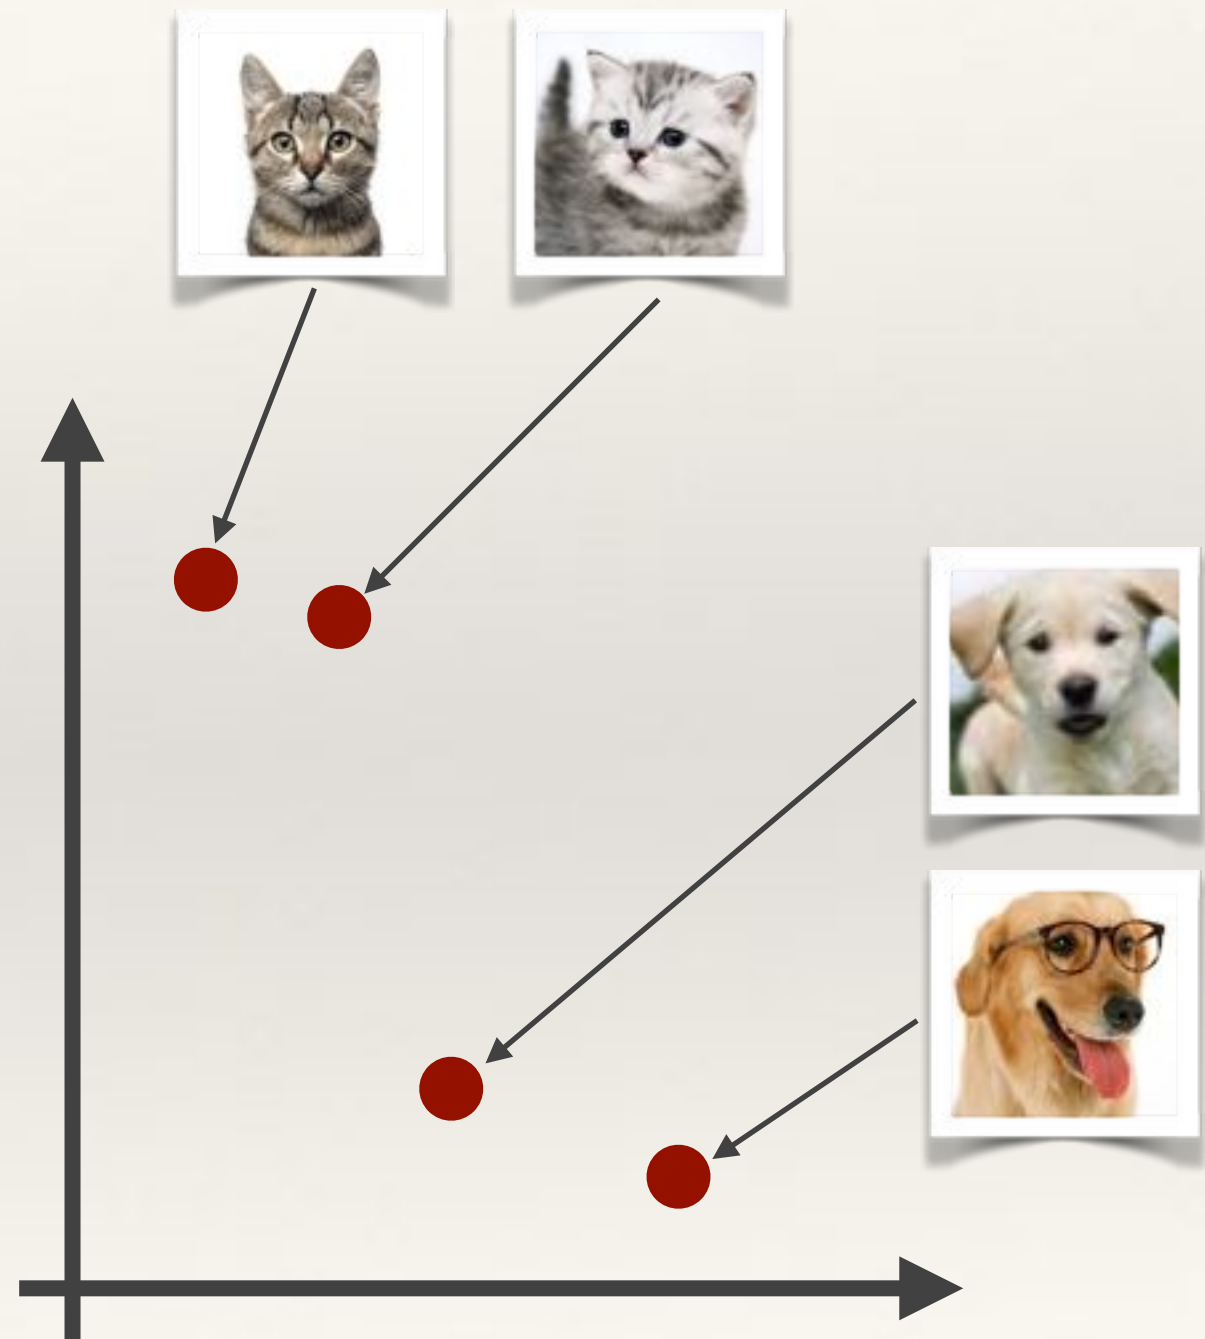
- ❖ **featurevector**: a mathematical vector
 - ❖ just a list of (usually Real) numbers
 - ❖ has a fixed number of **elements** in it
 - ❖ The number of elements is the **dimensionality** of the vector
- ❖ represents a **point** in a **featurespace** or equally a **direction** in the featurespace
- ❖ the **dimensionality of a featurespace** is the dimensionality of every vector within it
 - ❖ vectors of differing dimensionality can't exist in the same featurespace

*Demo: a really simple feature
extractor*

Distance and similarity

Distances in featurespace


- ❖ Feature extractors are often defined so that they produce vectors that are *close* together for *similar* inputs
- ❖ Closeness of two vectors can be computed in the feature space by measuring a distance between the vectors.



Euclidean distance (*L2 distance*)

- ❖ L2 distance is the most intuitive distance...
- ❖ The straight-line distance between two points
- ❖ Computed via an extension of Pythagoras theorem to n dimensions:


$$D_2(p, q) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2} = \|p - q\| = \sqrt{(p - q) \cdot (p - q)}$$



L1 distance (*aka Taxicab/Manhattan*)

- ❖ L1 distance is computed along paths parallel to the axes of the space:

$$D_1(p, q) = \sum_{i=1}^n |p_i - q_i| = \|p - q\|_1$$



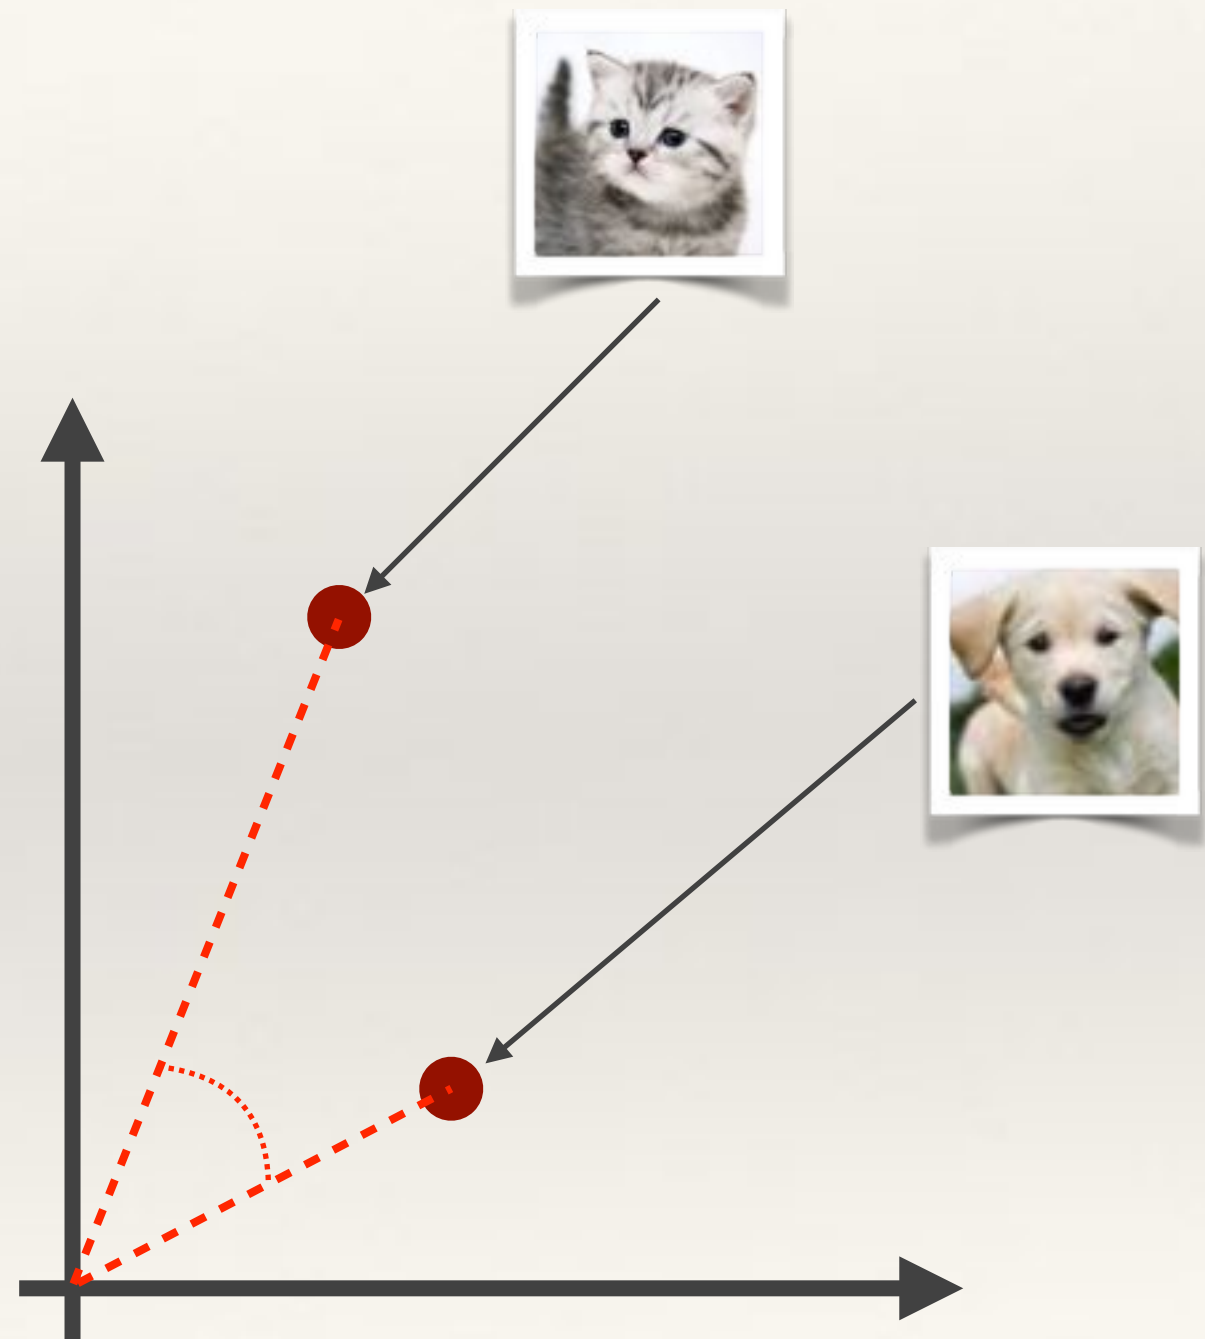
Cosine Similarity

- ❖ Cosine similarity measures the cosine of the angle between two vectors

- ❖ **It is not a distance!**

$$\cos(\theta) = \frac{p \cdot q}{\|p\| \|q\|} = \frac{\sum_{i=1}^n p_i q_i}{\sqrt{\sum_{i=1}^n p_i^2} \sqrt{\sum_{i=1}^n q_i^2}}$$

- ❖ Useful if you don't care about the relative length of the vectors



Choosing good feature vector representations for machine-learning

- ❖ Choose features which allow to distinguish objects or classes of interest
 - ❖ Similar within classes
 - ❖ Different between classes
- ❖ Keep number of features small
 - ❖ Machine-learning can get more difficult as dimensionality of featurespace gets large

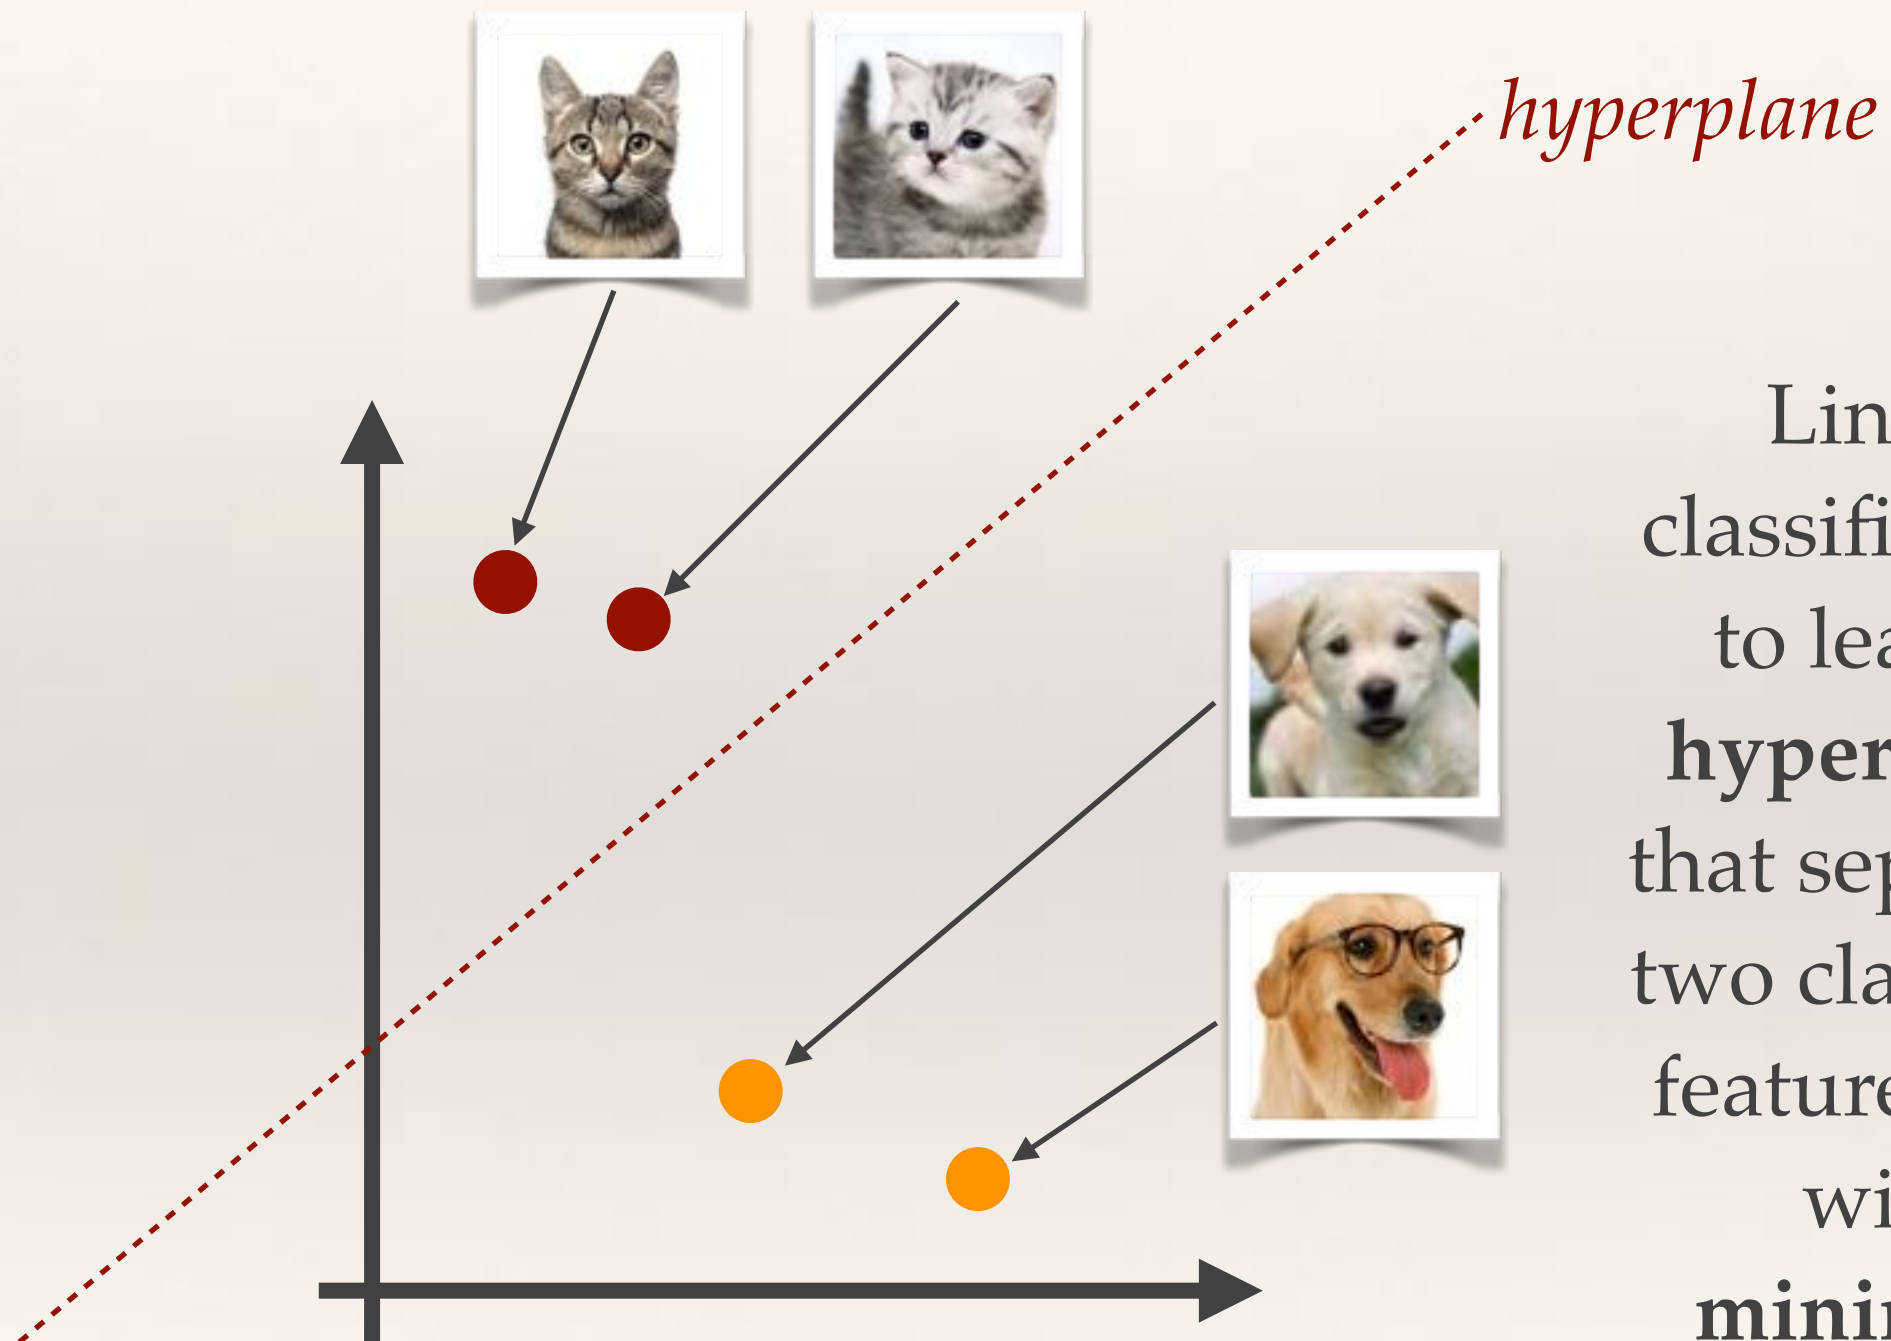
Supervised Machine Learning: *Classification*

- ❖ **Classification** is the process of assigning a **class label** to an object (typically represented by a vector in a feature space).
- ❖ A **supervised machine-learning algorithm** uses a set of pre-labelled *training data* to learn how to assign class labels to vectors (and the corresponding objects).
- ❖ A **binary** classifier only has two classes
- ❖ A **multiclass** classifier has many classes.



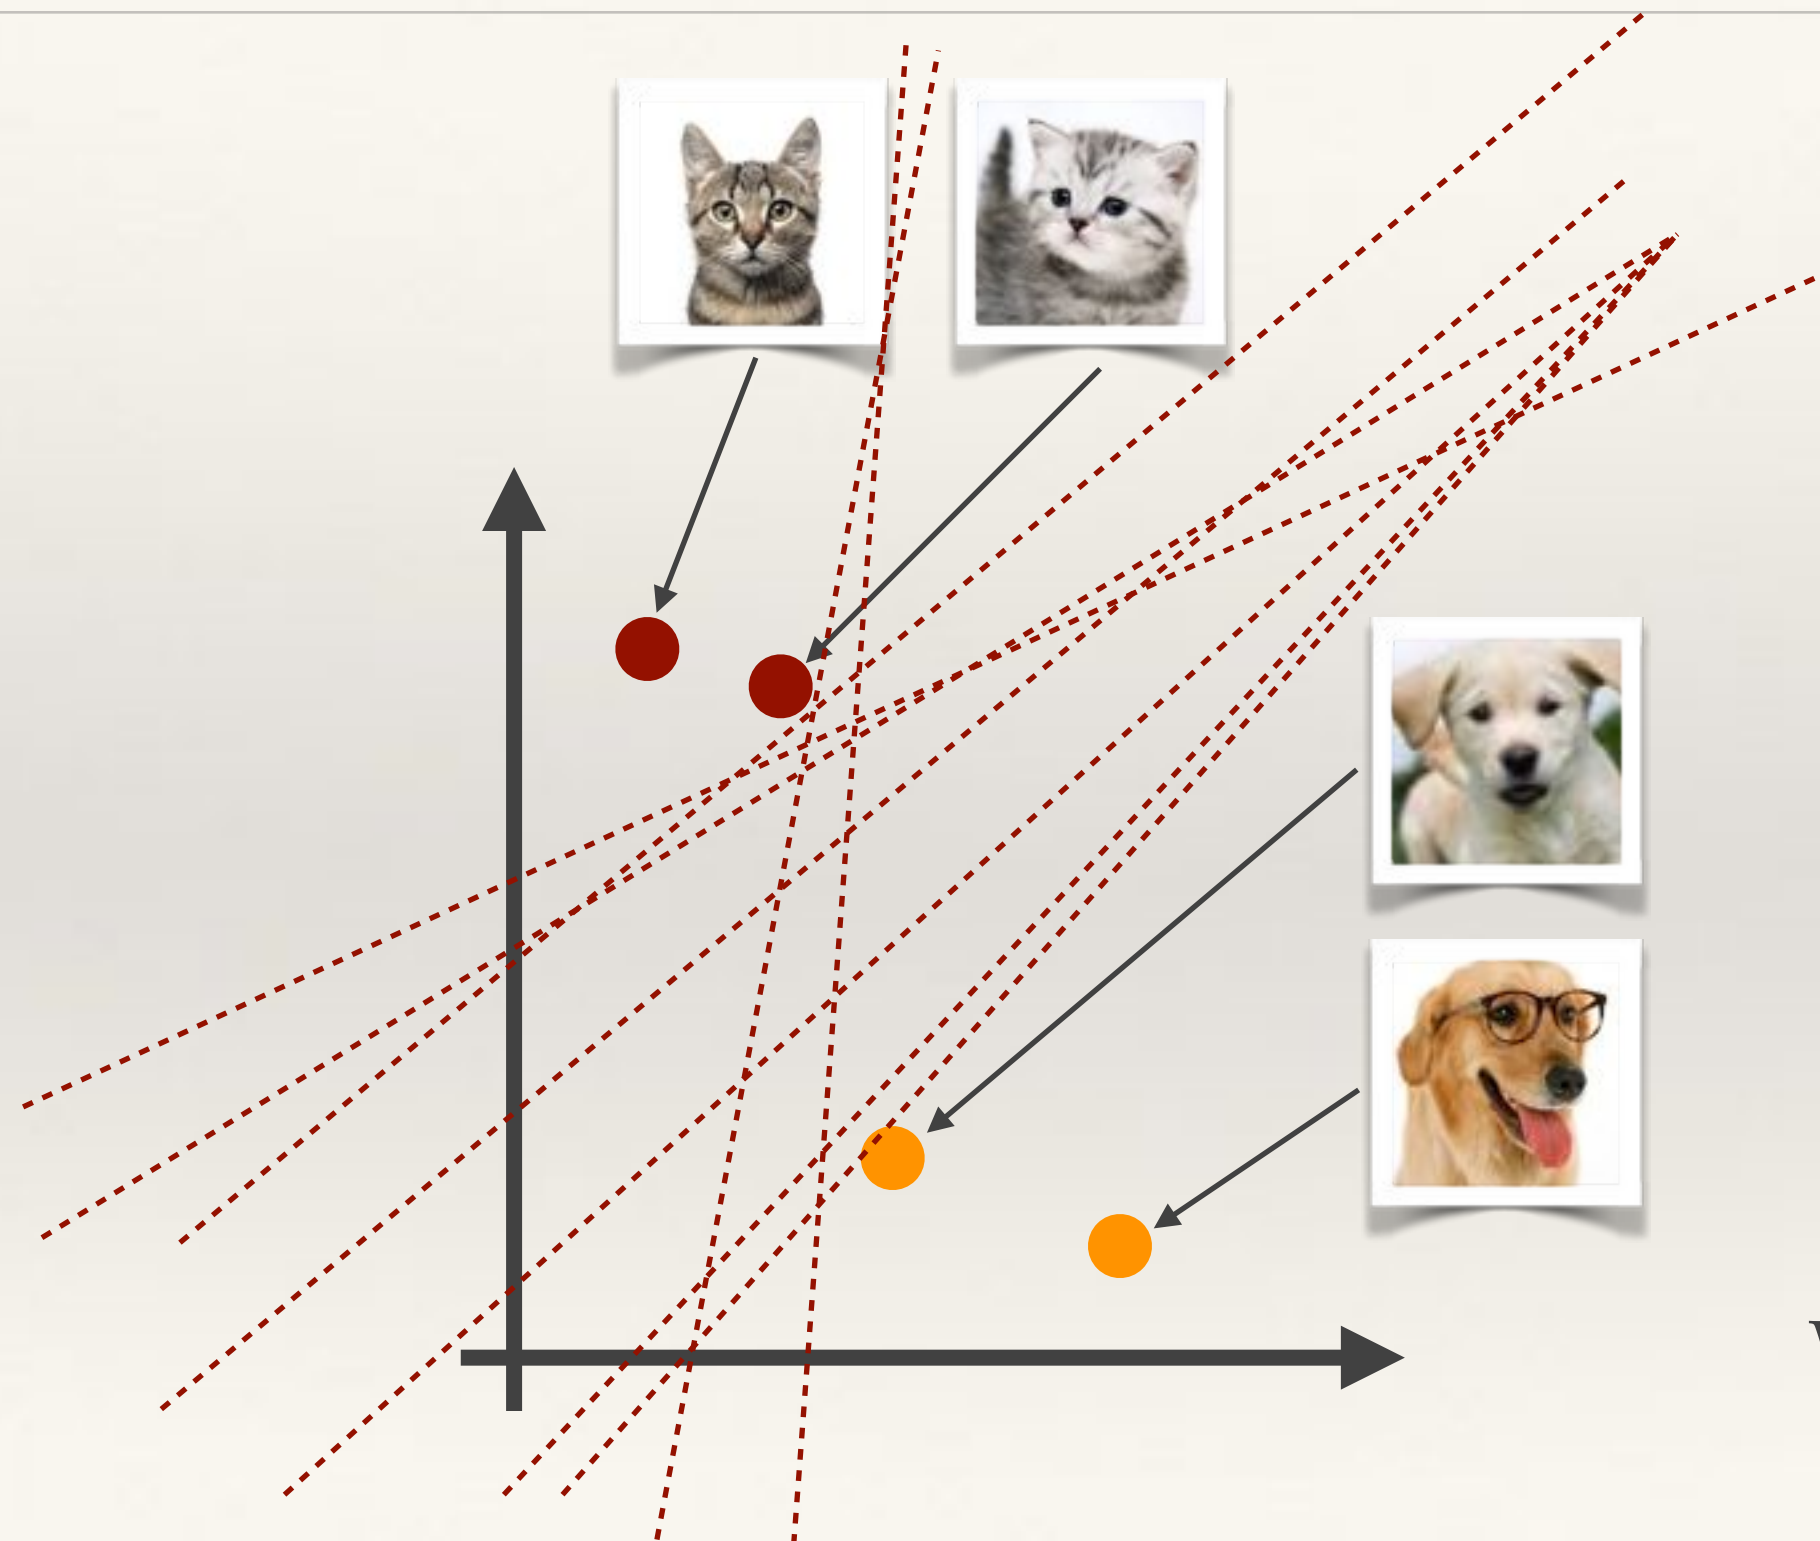
Cat or Dog?

Linear classifiers




Linear classifiers try to learn a **hyperplane** that separates two classes in featurespace with **minimum error**

Linear classifiers



Lots of
hyperplanes
to choose
from...
different
linear
classification
algorithms
apply
differing
constraints
when learning
the classifier


Linear classifiers



To classify a new image, you just need to check what side of the hyperplane it is on


perceptron demo

Non-linear binary classifiers




Linear
classifiers
work best
when the data
is linearly
separable...

Non-linear binary classifiers




No hope for a
linear
classifier!

Non-linear binary classifiers




Non-linear
binary
classifiers,
such as
**Kernel
Support
Vector
Machines**
learn non-
linear decision
boundaries

Non-linear binary classifiers



Have to be
careful... you
might lose
generality by
overfitting

Multiclass classifiers: KNN



... Must be a Cat!...

Assign class of
unknown
point based on
majority class
of *closest K*
neighbours in
featurespace

Demo: KNN Classification

KNN Problems

- ❖ Computationally expensive if there are:
 - ❖ Lots of training examples
 - ❖ Many dimensions

Unsupervised Machine Learning:

Clustering

- ❖ Clustering aims to group data without any prior knowledge of what the groups should look like or contain.
- ❖ In terms of feature vectors, items with similar vectors should be grouped together by a clustering operation.
- ❖ Some clustering operations create overlapping groups; for now we're only interested in disjoint clustering methods that assign an item to a single group.



K-Means Clustering

- ❖ K-Means is a classic featurespace clustering algorithm for grouping data into K groups with each group represented by a *centroid*:
 - ❖ The value of K is chosen
 - ❖ K initial cluster centres are chosen
 - ❖ Then the following process is performed iteratively until the centroids don't move between iterations:
 - ❖ Each point is assigned to its closest centroid
 - ❖ The centroid is recomputed as the mean of all the points assigned to it. If the centroid has no points assigned it is randomly re-initialised to a new point.
- ❖ The final clusters are created by assigning all points to their nearest centroid.

Demo: K-Means Clustering